


Statistical Learning Problem Set # 2

John Gaebler – joga5948@colorado.edu

February 10, 2020

1. Creation of Datasets

- (a) Write a function that generates samples - see appendix code
- (b) Display the true function $f(x)$ using 100 uniformly spaced samples on $[0,1]$ and the 30 noisy samples on the same plot.



The black points correspond to the true function, and the green points are the noisy data sample

2. Estimation of the Polynomial

- (a) Write a function that computes the optimal estimator \hat{y} given a fixed order p . Give the coefficients generated for $p = 1, 4$, and 15

Coefficients: $p = 1$

Intercept	x
1.0395355	-0.4407709

Coefficients: $p = 4$


Intercept	x	x^2	x^3	x^4
1.05401141	-1.24916710	2.23042888	-0.02250314	-1.71335123

Coefficients: $p = 15$

Intercept	x	x^2	x^3	x^4	-
-5.228157e-01	8.330017e+01	-1.528828e+03	1.263023e+04	-4.689995e+04	-
x^5	x^6	x^7	x^8	x^9	-
4.022791e+03	6.963936e+05	-3.157657e+06	7.470297e+06	-1.087213e+07	-
x^{10}	x^{11}	x^{12}	x^{13}	x^{14}	x^{15}
9.938791e+06	-5.401311e+06	1.415607e+06	NA	-5.829778e+04	NA

Note: The x^{13} and x^{15} coefficients in the $p = 15$ model are NA because they are linearly related to other components of the model. Past degree 14, NAs appear for the coefficients with very high frequency.

- (b) Display the three optimal polynomials as well as the true function and the noisy data on the same plot



Blue line: degree 1, Red line: degree 4, Pink line: degree 15

(c) Based on visual inspection, what is the optimal degree polynomial for:

i. The best fit to the noisy data?

Answer: The degree 15 polynomial best fits the noisy data points.

ii. The best fit to the true function?

Answer: The degree 4 polynomial best fits the path of the true function.

The higher degree polynomial can do a better job of fitting the variation in the data points, but it fails to capture the actual movement of the function. The lower degree polynomial may actually better approximate the true function even if it doesn't appear to follow the noisy data as well. We call the effect of matching the noisy data too closely and as a result losing the true function "over-fitting the model."

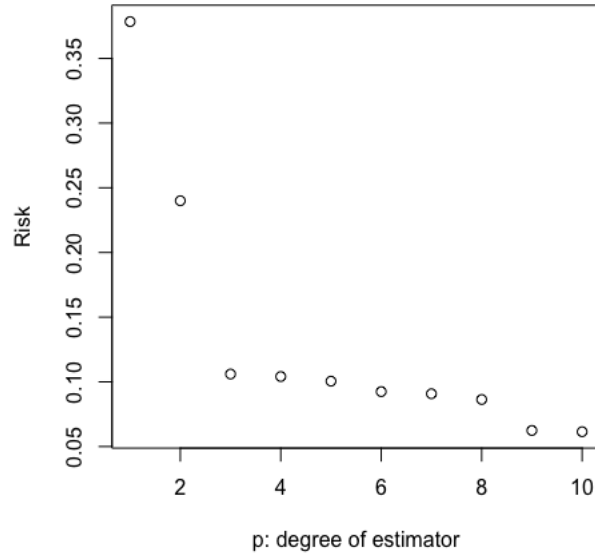
3. Choice of the Optimal Degree

- (a) For $p = 1, \dots, 10$ compute the values of \hat{R} . Display \hat{R} as a function of p , and find the value of p that minimizes \hat{R} .

Values of \hat{R} for each value of p :

$p = 1$	$p = 2$	$p = 3$	$p = 4$	$p = 5$
0.37845020	0.23997598	0.10603381	0.10410272	0.10051470
$p = 6$	$p = 7$	$p = 8$	$p = 9$	$p = 10$
0.09244547	0.09079719	0.08636445	0.06241269	0.06144891

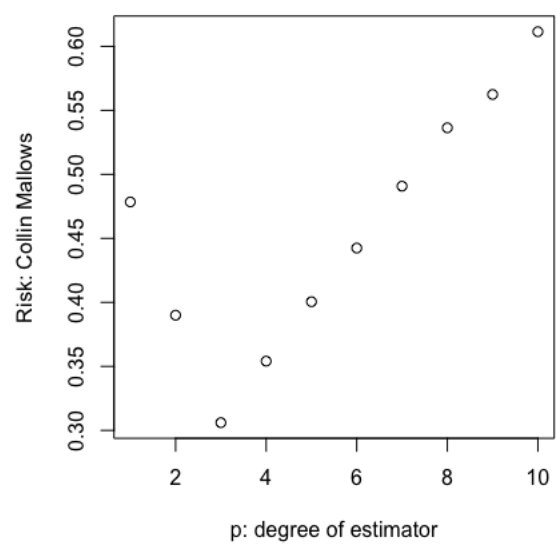
The value of p that minimizes \hat{R} is $p = 10$.



- (b) To penalize overly complex models, use the Collin Mallows statistic C_p to choose the optimal model. Display C_p as a function of p , and find the value of p that minimizes C_p .

$p = 1$	$p = 2$	$p = 3$	$p = 4$	$p = 5$
0.4784502	0.3899760	0.3060338	0.3541027	0.4005147
$p = 6$	$p = 7$	$p = 8$	$p = 9$	$p = 10$
0.4424455	0.4907972	0.5363644	0.5624127	0.6114489

The value of p that minimizes C_p is $p = 3$.




4. Estimate of the Bias and Variance of the Estimator


- (a) Let $S = 100$. Generate S random realizations of $y^{(s)}$ using the code that you wrote in the first problem. For each $y^{(s)}$, compute the optimal estimate $\hat{y}^{(s)}$. See appendix code. Finally, compute the Variance and the Bias of the estimator, and combine these statistics to compute the risk:

$$\tilde{R} = \text{Var}(\hat{y}) + \text{bias}^2(\hat{y})$$


for $p = 1, \dots, 10$. Then display $\text{Var}(\hat{y})$, $\text{bias}^2(\hat{y})$, and \tilde{R} as functions of p . Describe and explain the behavior of the curves as p increases. What is the value of p that minimizes \tilde{R} ?



The value of p that minimizes variance is $p = 10$




The value of p that minimizes bias squared is $p = 10$




The value of p that minimizes Risk is $p = 10$

For all of the statistics, the values decreases as p increases. This is because a function with higher degree is going to be able to better approximate the sample data. For Bias squared in particular, the value rapidly approaches zero as early as $p = 3$. However, the higher degree estimator leads to an overly complex model that is likely to over fit the data and produce predictions that are less accurate than a more simple model.

- (b) For each of the S data sets, compute the value p^* that minimizes the residual sum of squares. Additionally, compute the value p^+ that minimize the the statistic C_p .



For all the $S = 100$ samples, the degree 10 polynomial estimator produces the least value of Residual Sum of Squares



For all the $S = 100$ samples, the degree 3 polynomial estimator produces the least value of the Collin Mallows statistic.

The differences between these measures of risk demonstrates the importance of penalizing excessively complex models. There is an apparent large difference in the optimal degree of estimator between these two metrics of risk (3 vs 10). These estimates are very consistent across the $S = 100$ realizations of testing data, considering there is not a single deviation from $p = 10$ or $p = 3$ respectively.

Appendix: Code

```
1 ▾ #=====
2 ▾ ##-----Homework #2-----
3 library(ggplot2)
4 #Generating data
5 x_Data <- seq(1/30,1, by = 1/30)
6 true_fun_in <- runif(100, 0 ,1)
7
8 #function to generate the samples
9 ▾ generate_Data <- function(x, sd =0) {
10   data <- 1 - x + (2 * x ^2) - (.8 * x ^ 3) + (.06 * x ^4) - (x ^ 5)
11   data <- data + rnorm(30, 0, sd)
12   return(data)
13 }
14
15
16 y_Sample <- generate_Data(x_Data, .05)
17 true_fun_out <- generate_Data(true_fun_in)
18 both_in <- append(true_fun_in, x_Data)
19 both_out<- append(true_fun_out, y_Sample)
20 y_Data
21
22 bigplot <- ggplot()+
23   geom_point(aes(x = true_fun_in, y = true_fun_out))+
24   geom_point(aes(x = x_Data, y = y_Sample), color = "green")+
25   labs(x = "X Values", y = "Y Values")
26
27 ▾ ##-----
28 |
29 ▾ estimator_degree <- function(n, data = x_Data) {
30   i = 1
31   data_matrix <- array(data = rep(x_Data, n), dim = c(length(x_Data),n))
32   estimator <- c(rep(x_Data, n))
33 ▾ while (i <= n) {
34   data_matrix[,i] <- data_matrix[,i]^i
35   i = i +1
36 }
37 return(data_matrix)
38 }
39
40 ▾ #-----
41 p1_data <- estimator_degree(1)
42 p1 <- lm(y_Sample ~ p1_data)
43
44 degree1 <- function(x) 1.0395 + -0.4408 *x
45
```

```

44 degree1 <- function(x) 1.0395 + -0.4408 *x
45
46 p2_data <- estimator_degree(2)
47 p2 <- lm(y_Sample ~ p2_data)
48
49 p3_data <- estimator_degree(3)
50 p3 <- lm(y_Sample ~ p3_data)
51
52 p4_data <- estimator_degree(4)
53 p4 <- lm(y_Sample ~ p4_data)
54
55 p5_data <- estimator_degree(5)
56 p5 <- lm(y_Sample ~ p5_data)
57
58 p6_data <- estimator_degree(6)
59 p6 <- lm(y_Sample ~ p6_data)
60
61 p7_data <- estimator_degree(7)
62 p7 <- lm(y_Sample ~ p7_data)
63
64 p8_data <- estimator_degree(8)
65 p8 <- lm(y_Sample ~ p8_data)
66
67 p9_data <- estimator_degree(9)
68 p9 <- lm(y_Sample ~ p9_data)
69
70 p10_data <- estimator_degree(10)
71 p10 <- lm(y_Sample ~ p10_data)
72
73 data_vec <- c(p1_data, p2_data, p3_data, p4_data, p5_data, p6_data, p7_data, p8_data, p9_data, p10_data)
74
75
76
77 degree4 <- function(x) 1.05401 -1.24917*x +2.23043 *x^2 -0.02250 * x^3 -1.71335 *x^4
78
79 p15_data <- estimator_degree(15)
80 p15 <- lm(y_Sample ~ p15_data)
81
82 degree4 <- function(x) {
83   output <- 0
84   for(i in 1:length(coef(p4))) {
85     output <- output + as.numeric(coef(p4)[i]) * x ^(i-1)
86   }
87   return(output)
88 }

```

```

81
82 ▾ degree4 <- function(x) {
83   output <- 0
84 ▾   for(i in 1:length(coef(p4))) {
85     output <- output + as.numeric(coef(p4)[i]) * x ^ (i-1)
86   }
87   return(output)
88 }
89 ▾ degree15 <- function(x) {
90   output <- 0
91 ▾   for(i in 1:length(coef(p15))) {
92     output <- output + ifelse(is.na(as.numeric(coef(p15)[i]) * x ^ (i-1)), 0, as.numeric(coef(p15)[i]) * x ^ (i-1))
93   }
94   return(output)
95 }
96
97 ▾ degree_n <- function(x, degree) {
98   output <- 0
99   reg <- eval(parse(text = paste("p", degree, sep = "")))
100 ▾   for(i in 1:length(coef(reg))) {
101     output <- output + ifelse(is.na(as.numeric(coef(reg)[i]) * x ^ (i-1)), 0, as.numeric(coef(reg)[i]) * x ^ (i-1))
102   }
103   return(output)
104 }
105
106 bigplot + stat_function(fun = degree1, aes(x = x), data = data.frame(x = 0), color = "blue") +
107   stat_function(fun = degree4, aes(x = x), data = data.frame(x = 0), color = "red") +
108   stat_function(fun = degree15, aes(x = x), data = data.frame(x = 0), color = "pink") +
109   ylim(.2,1)
110
111
112
113 ▾ Risk <- function(deg) {
114   r_hat <- sum((y_Sample - degree_n(x_Data, deg))^2)
115   return(r_hat)
116 }
117
118 risk_vec <- sapply(c(1:10), Risk)
119 plot(x = 1:10, y = risk_vec,
120      xlab = "p: degree of estimator",
121      ylab = "Risk")
122
123
124 c_p <- risk_vec + 2 * (c(2:11)) * .025
125

```

```

123
124 c_p <- risk_vec + 2 * (c(2:11)) * .025
125
126 plot(x = 1:10, y = c_p,
127       xlab = "p: degree of estimator",
128       ylab = "Risk: Collin Mallows")
129 ##----- Looks like the optimal degree is 3
130
131
132 ##-- Creatin the S = 100 random realization of y(s)
133 y_s <- rerun(100, generate_Data(x_Data, .05))
134
135 model_creator <- function(y, data) {
136   return(lm(y ~ data))
137 }
138 ## generating the p = {1,...10} models for each of the S = 100 realizations, 1000 total models
139 y_s_model_p1 <- lapply(y_s, model_creator, data = p1_data)
140 y_s_model_p2 <- lapply(y_s, model_creator, data = p2_data)
141 y_s_model_p3 <- lapply(y_s, model_creator, data = p3_data)
142 y_s_model_p4 <- lapply(y_s, model_creator, data = p4_data)
143 y_s_model_p5 <- lapply(y_s, model_creator, data = p5_data)
144 y_s_model_p6 <- lapply(y_s, model_creator, data = p6_data)
145 y_s_model_p7 <- lapply(y_s, model_creator, data = p7_data)
146 y_s_model_p8 <- lapply(y_s, model_creator, data = p8_data)
147 y_s_model_p9 <- lapply(y_s, model_creator, data = p9_data)
148 y_s_model_p10 <- lapply(y_s, model_creator, data = p10_data)
149
150 model_vec <- list(y_s_model_p1, y_s_model_p2, y_s_model_p3, y_s_model_p4, y_s_model_p5, y_s_model_p6,
151                  y_s_model_p7, y_s_model_p8, y_s_model_p9, y_s_model_p10)
152
153 mass_predictor <- function(model) {
154   fitted <- 0
155   reg <- model
156   for(i in 1:length(coef(reg))) {
157     fitted <- fitted + ifelse(is.na(as.numeric(coef(reg)[i])) * x_Data ^ (i-1)), 0,
158                               as.numeric(coef(reg)[i]) * x_Data ^ (i-1))
159   }
160   return(fitted)
161 }
162 ## creating fitted data with the models, 10 fitted sets per S = 100 datasets, 1000 fitted datasets
163 fitted_vec_p1 <- lapply(y_s_model_p1, mass_predictor)
164 fitted_vec_p2 <- lapply(y_s_model_p2, mass_predictor)
165 fitted_vec_p3 <- lapply(y_s_model_p3, mass_predictor)
166 fitted_vec_p4 <- lapply(y_s_model_p4, mass_predictor)
167 fitted_vec_p5 <- lapply(y_s_model_p5, mass_predictor)

```

```

162 ## creating fitted data with the models, 10 fitted sets per S = 100 datasets, 1000 fitted datasets
163 fitted_vec_p1 <- lapply(y_s_model_p1, mass_predictor)
164 fitted_vec_p2 <- lapply(y_s_model_p2, mass_predictor)
165 fitted_vec_p3 <- lapply(y_s_model_p3, mass_predictor)
166 fitted_vec_p4 <- lapply(y_s_model_p4, mass_predictor)
167 fitted_vec_p5 <- lapply(y_s_model_p5, mass_predictor)
168 fitted_vec_p6 <- lapply(y_s_model_p6, mass_predictor)
169 fitted_vec_p7 <- lapply(y_s_model_p7, mass_predictor)
170 fitted_vec_p8 <- lapply(y_s_model_p8, mass_predictor)
171 fitted_vec_p9 <- lapply(y_s_model_p9, mass_predictor)
172 fitted_vec_p10 <- lapply(y_s_model_p10, mass_predictor)
173
174 fitted_vec <- list(fitted_vec_p1, fitted_vec_p2, fitted_vec_p3, fitted_vec_p4, fitted_vec_p5,
175                  fitted_vec_p6, fitted_vec_p7, fitted_vec_p8, fitted_vec_p9, fitted_vec_p10)
176
177
178 Var_calc <- function(fit) {
179   b <- mapply('-', y_s, fit)
180   RSS <- 0
181   for(i in 1:100) {
182     RSS <- RSS + sqrt(sum(b[,i]^2))^2
183   }
184   return(RSS / 99)
185 }
186
187 Variance_vec <- lapply(fitted_vec, Var_calc)
188 # creating the sample means for the optimal estimated for each p = {1, ... , 10}
189 y_bar <- lapply(fitted_vec, function(x) Reduce("+", x) / 99)
190
191 true_Values <- generate_Data(x_Data)
192 true_Values_vec <- list(true_Values, true_Values, true_Values, true_Values, true_Values,
193                        true_Values, true_Values, true_Values, true_Values, true_Values)
194
195 bias <- mapply('-', true_Values_vec, y_bar)
196 for(i in 1:10) {
197   bias_squared[i] <- sqrt(sum(bias[,i]^2))^2
198 }
199
200
201 Risk_Model <- mapply('+', Variance_vec, bias_squared)
202

```

```

203 ## Plots
204
205 plot(x = 1:10, y = Variance_vec,
206      xlab = "p: degree of estimator",
207      ylab = "Variance of Estimator")
208
209 plot(x = 1:10, y = bias_squared,
210      xlab = "p: degree of estimator",
211      ylab = "Bias Squared of Estimator")
212
213 plot(x = 1:10, y = Risk_Model,
214      xlab = "p: degree of estimator",
215      ylab = "Risk of Estimator")
216
217 match(min(Risk_Model), Risk_Model)
218
219 ## compute the RSS for the models p 1-10
220
221 RSS <- function( model) {
222   output <- mapply('-', y_s, model)
223   output1 <- c(1:100)
224   for(i in 1:100) {
225     output[,i] <- output[,i]^2
226     output1[i] <- sum(output[,i])
227   }
228   return(output1)
229 }
230
231 p_RSS <- lapply(fitted_vec, RSS)
232
233 p_star <- c(1:100)
234
235 for(i in 1:100) {
236   vec <- c(1:10)
237   for(j in 1:10) {
238     vec[j] <- p_RSS[[j]][i]
239   }
240   p_star[i] <- match(min(vec), vec)
241 }
242

```

```

219 ## compute the RSS for the models p 1-10
220
221 RSS <- function( model) {
222   output <- mapply('-', y_s, model)
223   output1 <- c(1:100)
224   for(i in 1:100) {
225     output[,i] <- output[,i]^2
226     output1[i] <- sum(output[,i])
227   }
228   return(output1)
229 }
230
231 p_RSS <- lapply(fitted_vec, RSS)
232
233 p_star <- c(1:100)
234
235 for(i in 1:100) {
236   vec <- c(1:10)
237   for(j in 1:10) {
238     vec[j] <- p_RSS[[j]][i]
239   }
240   p_star[i] <- match(min(vec), vec)
241 }
242
243
244 plot(x = 1:100, y = p_star,
245      ylab = "p* value for lowest RSS",
246      xlab = "S = {1, ... , 100} samples")
247
248
249
250 c_of_p <- c(1:100)
251
252 for(i in 1:100) {
253   vec <- c(1:10)
254   for(j in 1:10) {
255     vec[j] <- p_RSS[[j]][i] + 2* Variance_vec[[j]] * (j +1)
256   }
257   c_of_p[i] <- match(min(vec), vec)
258 }
259
260 plot(x = 1:100, y = c_of_p,
261      ylab = "p-dagger value for lowest Collin Mallows Statistic",
262      xlab = "S = {1, ... , 100} samples")
263

```